

Qwen3-VL-30B-A3B-Thinking: Best AI For Long Video Analysis: Run Locally



Fahd Mirza
47.4K subscribers

Join

Subscribe

4



Share

Download

Thanks



28 views Oct 5, 2025 #qwen3vl

This video locally installs Qwen3-VL-30B-A3B-Thinking, which handles books and hours-long video with full recall and second-level indexing.

Hugging Face

Search models, datasets, Models Datasets Spaces Docs Pricing Log In Sign Up

Qwen/Qwen3-VL-30B-A3B-Thinking like 49 Follow Qwen 52.6k

Safetensors qwen3_vl_moe arxiv:4 papers License: apache-2.0

Model card Files xet Community 1

Qwen Chat

Qwen3-VL-30B-A3B-Thinking

Meet Qwen3-VL — the most powerful vision-language model in the Qwen series to date.

This generation delivers comprehensive upgrades across the board: superior text understanding & generation, deeper visual perception & reasoning, extended context length, enhanced

Downloads last month 88

Safetensors ①

Model size 31.1B params Tensor type BF16

Chat template Files info

Inference Providers NEW

This model runs with about 65GB of vRAM

Home Videos Shorts Live Podcasts Playlists qwen3-VL

MIND BLOWING QWEN3-VL 30B 10:11

Qwen3-VL — The Most Powerful Vision-Language Model: Install and Test

Fahd Mirza • 313 views • 2 hours ago

This video locally installs Qwen3-VL — the most powerful vision-language model in the Qwen series to date. 🍌


Buy Me a Coffee to support the channel: <https://ko-fi.com/fahdmirza> #qwen3vl...

QWEN3-NEXT NO GPU NEEDED! RUN QUANT LOCALLY 8:14

Qwen3-Next 80B Int4 Quant on CPU — No GPU Needed! Easy Tutorial

Fahd Mirza • 9.1K views • 2 weeks ago

This video installs and test Qwen3-Next-80B-A3B-Thinking-int4-mixed-AutoRound locally. 🍌 This video is sponsored by <https://www.eigent.ai>, the World's First Multi-agent Workforce to Unlock...



		Qwen3-VL 30B-A3B Instruct	Qwen3-30B-A3B Instruct-2507	Qwen3-VL 30B-A3B Thinking	Qwen3-30B-A3B Thinking-2507
Knowledge	MMLU	85.0	85.0	87.6	87.3
	MMLU-Pro	77.8	78.4	80.5	80.9
	MMLU-Redux	88.4	89.3	90.9	91.4
	GPQA	70.4	70.4	74.4	73.4
	SuperGPQA	53.1	53.4	56.4	56.8
	SimpleQA	27.0	22.2	23.9	19.2
Reasoning	AIME25	69.3	61.3	83.1	85.0
	HMMT25	50.6	43.0	67.6	71.4
	LiveBench1125	65.4	69.0	72.1	76.8
Code	LCBv6 (25.02 - 25.05)	42.6	43.2	64.2	66.0
Instruction Following	SIFO	50.1	46.8	66.9	66.9
	SIFO-multiturn	35.1	36.4	60.3	59.3
	IFEval	85.8	84.7	81.7	88.9
Subjective Evaluation	Arena-Hard v2	58.5	69.0	56.7	56.0
	Creative Writing v3	84.6	86.0	82.5	84.4
	WritingBench	82.6	85.5	85.2	85.0
Agent	BFCL-v3	66.3	65.1	68.6	72.4
	MultitF	66.1	67.9	73.0	76.4
Multilingual	MMLU-ProX	70.9	72.0	76.1	76.4
	INCLUDE	71.6	71.9	74.5	74.4
	PolyMATH	44.3	43.1	51.7	52.6




For reproducibility of Arena-Hard v2, we report the win rates evaluated by GPT-4.1.



```

https://huggingface.co/Qwen/Qwen3-VL-30B-A3B-Thinking
170%
Ubuntu@0151-dsm-prmx30007: ~
(ai) Ubuntu@0151-dsm-prmx30007:~$ pip install git+https://github.com/huggingface/transformers

```

```
Ubuntu@0151-dsm-prxm30007:~$ jupyter notebook
[I 2025-10-05 01:07:23.075 ServerApp] jupyter_lsp | extension was successfully linked.
[I 2025-10-05 01:07:23.079 ServerApp] jupyter_server_terminals | extension was successfully linked.
[I 2025-10-05 01:07:23.082 ServerApp] jupyterlab | extension was successfully linked.
[I 2025-10-05 01:07:23.085 ServerApp] notebook | extension was successfully linked.
```





localhost:8888/notebooks/Untitled1.ipynb 150%

Jupyter Untitled1 Last Checkpoint: 2 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
merges.txt: 1.67M? [00:00<00:00, 8.52MB/s]
tokenizer.json: 7.03M? [00:00<00:00, 20.2MB/s]
video_preprocessor_config.json: 100% 385/385 [00:00<00:00, 62.9kB/s]
chat_template.json: 5.41k? [00:00<00:00, 782kB/s]
```

```
[ ]: vid_path = "/home/Ubuntu/myvideos/couplewalking.mp4"
prompt="Describe the video briefly and note any key actions."

messages = [{"role": "user", "content": [{"type": "video", "video": vid_path}, {"type": "text", "text": prompt}],}]


inputs = processor.apply_chat_template( messages, tokenize=True, add_generation_prompt=True, return_dict=True, return_tensors="pt")

device = next(model.parameters()).device
inputs = {k: (v.to(device) if isinstance(v, torch.Tensor) else v) for k, v in inputs.items()}

if getattr(model.generation_config, "pad_token_id", None) is None:
    model.generation_config.pad_token_id = model.generation_config.eos_token_id
model.eval()

with torch.inference_mode():
    generated_ids = model.generate(**inputs, max_new_tokens=512, use_cache=True)

new_tokens = generated_ids[:, inputs["input_ids"].shape[1]:]
text = processor.batch_decode(new_tokens, skip_special_tokens=True, clean_up_tokenization_spaces=False)
```



The screenshot shows a JupyterLab window titled 'Untitled1' with a last checkpoint of 2 minutes ago. The code cell contains the following Python code:

```
messages = [{"role": "user", "content": [{"type": "video", "video": vid_path}, {"type": "text", "text": prompt}],}]

inputs = processor.apply_chat_template( messages, tokenize=True, add_generation_prompt=True, return_dict=True, return_tensors="pt")


device = next(model.parameters()).device
inputs = {k: (v.to(device) if isinstance(v, torch.Tensor) else v) for k, v in inputs.items()}

if getattr(model.generation_config, "pad_token_id", None) is None:
    model.generation_config.pad_token_id = model.generation_config.eos_token_id
model.eval()

with torch.inference_mode():
    generated_ids = model.generate(**inputs, max_new_tokens=512, use_cache=True)

new_tokens = generated_ids[:, inputs["input_ids"].shape[1]:]
text = processor.batch_decode(new_tokens, skip_special_tokens=True, clean_up_tokenization_spaces=False)

display(Video(vid_path))
print(text[0])
```




The screenshot shows a JupyterLab window titled 'Untitled6' with a last checkpoint of 1 minute ago. The code cell contains the following Python code:

```
[2]: from transformers import Qwen3VLMOForConditionalGeneration, AutoProcessor
import torch
from IPython.display import Video

model = Qwen3VLMOForConditionalGeneration.from_pretrained("Qwen/Qwen3-VL-30B-A3B-Thinking", dtype=torch.bfloat16, attn_implementation="flash")
processor = AutoProcessor.from_pretrained("Qwen/Qwen3-VL-30B-A3B-Thinking")
```

Below the code cell, a progress bar indicates the loading of checkpoint shards: 100% (13/13) [00:12<00:00, 1.15it/s].

The model is now downloaded and we can proceed with the inferencing as below



The screenshot shows a JupyterLab window titled 'Untitled6' with a last checkpoint of 2 minutes ago. The code cell contains the following Python code:

```
[ ]: vid_path = "/home/Ubuntu/myvideos/couplewalking.mp4"
prompt="Describe the video briefly and note any key actions."

messages = [{"role": "user", "content": [{"type": "video", "video": vid_path}, {"type": "text", "text": prompt}],}]

inputs = processor.apply_chat_template( messages, tokenize=True, add_generation_prompt=True, return_dict=True, return_tensors="pt")


device = next(model.parameters()).device
inputs = {k: (v.to(device) if isinstance(v, torch.Tensor) else v) for k, v in inputs.items()}

if getattr(model.generation_config, "pad_token_id", None) is None:
    model.generation_config.pad_token_id = model.generation_config.eos_token_id
model.eval()

with torch.inference_mode():
    generated_ids = model.generate(**inputs, max_new_tokens=512, use_cache=True)

new_tokens = generated_ids[:, inputs["input_ids"].shape[1]:]
text = processor.batch_decode(new_tokens, skip_special_tokens=True, clean_up_tokenization_spaces=False)

display(Video(vid_path))
print(text[0])
```



The screenshot shows a JupyterLab environment. On the left, a video player titled 'angry.mp4 - VLC media player' displays a man with a frustrated expression, holding his hands to his head. The video file 'angry.mp4' is listed in the left sidebar. The main area is a code editor with a dark theme, showing Python code for a video generation model. The code includes imports for torch, torchvision, and transformers, and defines a function 'generate' that takes inputs and returns a video path. The code is partially obscured by the video player window.


```
localhost:8888/notebooks/Untitled6.ipynb 150%
Jupyter Untitled6 Last Checkpoint: 14 minutes ago
File Edit View Run Kernel Settings Help Trusted
JupyterLab Python 3 (ipykernel)

In [8]: vid_path = "/home/ubuntu/myvideos/angry.mp4"
prompt = """
What are the emotions of this person in this video?
And what could be the reason for those emotions?
One hint is that this man always asks the viewers to like, subscribe and become member of channel.
Could that be the reason? I
"""

messages = [{"role": "user", "content": [{"type": "video", "video": vid_path}, {"type": "text", "text": prompt}],}]

inputs = processor.apply_chat_template(messages, tokenize=True, add_generation_prompt=True, return_dict=True, return_tensors="pt")

device = next(model.parameters()).device
inputs = {k: (v.to(device) if isinstance(v, torch.Tensor) else v) for k, v in inputs.items()}

if getattr(model.generation_config, "pad_token_id", None) is None:
    model.generation_config.pad_token_id = model.generation_config.eos_token_id
model.eval()

with torch.inference_mode():
    generated_ids = model.generate(**inputs, max_new_tokens=4512, use_cache=True)

new_tokens = generated_ids[:, inputs["input_ids"].shape[1]:]
text = processor.batch_decode(new_tokens, skip_special_tokens=True, clean_up_tokenization_spaces=False)

print(text[0])
```

```
localhost:8888/notebooks/Untitled6.ipynb 150%
Jupyter Untitled6 Last Checkpoint: 15 minutes ago
File Edit View Run Kernel Settings Help Trusted
JupyterLab Python 3 (ipykernel)

print(text[0])

So, let's analyze this. First, the person's expression: he looks stressed or frustrated. His eyes are closed in the second frame, maybe sighing or showing exasperation. The hint says he's always asking for likes, subs, members. So maybe he's frustrated with not getting enough engagement? Like, the channel isn't growing as much as he wants, so he's stressed about the audience not subscribing or liking. That makes sense. People who do this often feel pressure to get more viewers, and if they're not getting the support they need, it leads to stress or frustration. So emotions are likely stress, frustration, maybe anxiety. Reason: possibly not getting enough likes/subscriptions, which is a common issue for content creators trying to grow their channel. The hint points to that, so that's the likely reason.

</think>

The person in the video appears to be experiencing **stress, frustration, or exasperation** (evidenced by their facial expression—such as closed eyes in the second frame, which can indicate sighing or emotional strain).

Given the hint that this individual consistently asks viewers to like, subscribe, and become channel members, the most plausible reason for these emotions is **frustration related to viewer engagement**. Content creators often feel pressure to grow their audience, and if they perceive a lack of likes, subscriptions, or member support (e.g., low engagement with their calls-to-action), it can lead to stress, anxiety, or frustration about their channel's growth or success. This is a common emotional response for creators who rely on audience interaction to sustain or expand their platform.

In short: The emotions likely stem from *disappointment or stress about not receiving sufficient engagement* (likes, subscriptions, member sign-ups) from viewers, which is critical for many content creators' goals.
```