

WanGP - AI Video Generation for GPU Poor - Use Locally with Any Model



Fahd Mirza
47.6K subscribers

Join

Subscribe

11



Share

Ask

Download



156 views Oct 9, 2025 #wangp #wan2gp

This video is a step-by-step Tutorial to install and use Wan2GP. It supports the Wan (and derived models), Hunyuan Video and LTV Video models.

https://github.com/deepbeepmeep/Wan2GP

README License

WanGP

WanGP by DeepBeepMeep : The best Open Source Video Generative Models Accessible to the GPU Poor

WanGP supports the Wan (and derived models), Hunyuan Video and LTV Video models with:

- Low VRAM requirements (as low as 6 GB of VRAM is sufficient for certain models)
- Support for old GPUs (RTX 10XX, 20xx, ...)
- Very Fast on the latest GPUs
- Easy to use Full Web based interface
- Auto download of the required model adapted to your specific architecture
- Tools integrated to facilitate Video Generation : Mask Editor, Prompt Enhancer, Temporal and Spatial Generation, MMAudio, Video Browser, Pose / Depth / Flow...

WAN2.1

Pushes Limits of AI Video Generation



Install & Test Locally 10:21

WAN Video Model from Alibaba - Install Locally for Text to Video - Hands-on

Fahd Mirza • 9.5K views • 7 months ago

This video shows how to locally install Wan2.1 AI model from Alibaba which is designed on the mainstream diffusion transformer paradigm. 🧨 Get 50% Discount on any A6000 or A5000 GPU rental...



Install & Test Locally 14:32

Install Wan2.2 TI2V 5B Locally and Test for Text-to-Video and Image-to-Video

Fahd Mirza • 6.7K views • 2 months ago

This video locally installs Wan2.2, which is the premier video foundation model. 🧨 This video is sponsored by <https://www.eigent.ai>, the World's First Multi-agent Workforce to Unlock Your...



Tencent Releases Hunyuan3D World Model 1.0 - Install Locally - Free and Open-Source

Fahd Mirza • 5.1K views • 2 months ago

This video locally installs Hunyuan3D World Model 1.0 which enables you to interactive 3D worlds from text and image conditions. 🎁 Get 50% Discount on any A6000 or A5000 GPU rental, use...



Hunyuan-7B: Qwen3 Rival Is Here: Install and Test Locally for Reasoning Magic

Fahd Mirza • 3K views • 2 months ago

This video locally installs Hunyuan-7B-instruct model from Tencent. 🎁 This video is sponsored by <https://www.eigant.ai>, the World's First Multi-agent Workforce to Unlock Your Exceptional...

```

CN+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
TE Ubuntu@0360-kci2-ty6k-prxmx100037:~$ conda create -n ai python=3.11 -y && conda activate ai
Channels:
Blk: - defaults
Sh Platform: linux-64
a1bu Collecting package metadata (repodata.json): /

```

```

W Downloading and Extracting Packages:
Lk
Yo Preparing transaction: done
Verifying transaction: done
Executing transaction: done
bscrip#
# To activate this environment, use
#
# $ conda activate ai
#
# To deactivate an active environment, use
#
# $ conda deactivate
Sh
(ai) Ubuntu@0360-kci2-ty6k-prxmx100037:~$
a1bu

```

```


(ai) Ubuntu@0360-kci2-ty6k-prxmx100037:~$ git clone https://github.com/deepbeepmeep/Wan2GP.git && cd Wan2GP
Cloning into 'Wan2GP'...
remote: Enumerating objects: 4330, done.
remote: Counting objects: 100% (543/543), done.
remote: Compressing objects: 100% (266/266), done.
remote: Total 4330 (delta 387), reused 315 (delta 276), pack-reused 3787 (from 2)
Receiving objects: 100% (4330/4330), 10.47 MiB | 27.98 MiB/s, done.
Resolving deltas: 100% (2713/2713), done.
(ai) Ubuntu@0360-kci2-ty6k-prxmx100037:~/Wan2GP$

```


```



(ai) Ubuntu@0360-kci2-ty6k-prxmx100037:~/Wan2GP$ docker --version
Docker version 26.1.0, build 9714adc
(ai) Ubuntu@0360-kci2-ty6k-prxmx100037:~/Wan2GP$


```

The application is now running on **localhost:7860** as above







```

Common Issue Diagnostics:
⚠ LD_LIBRARY_PATH may not include CUDA libraries

Starting application...

* Running on local URL: http://0.0.0.0:7860
* To create a public link, set `share=True` in `launch()`.

pose/dw-ll_ucoco_384.onnx: 100%| 134M/134M [00:01<00:00, 67.6MB/s]
pose/yolox_l.onnx: 100%| 217M/217M [00:01<00:00, 121MB/s]
scribble/net6_A_latest.pth: 100%| 17.2M/17.2M [00:00<00:00, 36.7MB/s]
flow/raft-things.pth: 100%| 21.1M/21.1M [00:00<00:00, 47.3MB/s]
depth/depth_anything_v2_vitl.pth: 100%| 1.34G/1.34G [00:03<00:00, 386MB/s]
depth/depth_anything_v2_vitb.pth: 100%| 390M/390M [00:02<00:00, 185MB/s]
mask/sam_vit_h_4b8939_fp16.safetensors: 100%| 1.28G/1.28G [00:03<00:00, 378MB/s]
mask/model.safetensors: 100%| 141M/141M [00:01<00:00, 102MB/s]
config.json: 2.90kB [00:00, 12.9MB/s]
config.json: 1.60kB [00:00, 8.77MB/s]
feature_extractor_config.json: 1.60kB [00:00, 9.70MB/s]
wav2vec/model.safetensors: 100%| 378M/378M [00:01<00:00, 248MB/s]
preprocessor_config.json: 100%| 159/159 [00:00<00:00, 1.44MB/s]
special_tokens_map.json: 100%| 85.0/85.0 [00:00<00:00, 803kB/s]
tokenizer_config.json: 100%| 163/163 [00:00<00:00, 1.51MB/s]
vocab.json: 100%| 291/291 [00:00<00:00, 2.09MB/s]
config.json: 1.95kB [00:00, 8.64MB/s]
chinese-wav2vec2-base/pytorch_model.bin: 100%| 380M/380M [00:01<00:00, 284MB/s]
preprocessor_config.json: 100%| 160/160 [00:00<00:00, 1.12MB/s]
roformer/model_bs_roformer_ep_317_sdr_12(...): 4%| 27.9M/639M [00:00<00:14, 42.8MB/s]

```

It then downloads the model for the first time as above before running the steps

```
Ubuntu@0360-kcl2-ty6k-prxm100037: ~/Wan2GP
flownet.pkl: 100%| 12.2M/12.2M [00:00<00:00, 37.5MB/s]
wan2.1_text2video_1.3B_mbf16.safetensors: 100%| 2.87G/2.87G [00:05<00:00, 486MB/s]
xlm-roberta-large/models_clip_open-clip-(...): 100%| 2.39G/2.39G [00:05<00:00, 477MB/s]
xlm-roberta-large/sentencepiece.bpe.model(...): 100%| 5.07M/5.07M [00:00<00:00, 17.5MB/s]
special_tokens_map.json: 100%| 280/280 [00:00<00:00, 2.28MB/s]
xlm-roberta-large/tokenizer.json: 100%| 17.1M/17.1M [00:00<00:00, 26.5MB/s]
tokenizer_config.json: 100%| 418/418 [00:00<00:00, 2.31MB/s]
special_tokens_map.json: 6.62kB [00:00, 20.5MB/s]
umt5-xxl/spiece.model: 100%| 4.55M/4.55M [00:00<00:00, 11.4MB/s]
umt5-xxl/tokenizer.json: 100%| 16.8M/16.8M [00:00<00:00, 40.9MB/s]
tokenizer_config.json: 61.7kB [00:00, 98.7MB/s]
umt5-xxl/models_t5 umt5-xxl-enc-quanto i(...): 100%| 6.73G/6.73G [00:13<00:00, 506MB/s]
Wan2.1_VAE.safetensors: 100%| 508M/508M [00:02<00:00, 234MB/s]
fantasy_proj_model.safetensors: 100%| 6.31M/6.31M [00:00<00:00, 14.8MB/s]
Loading Model 'ckpts/wan2.1_text2video_1.3B_mbf16.safetensors' ...
***** Memory Management for the GPU Poor (mmgp 3.6.2) by DeepBeepMeep *****
Pinning data of 'transformer' to reserved RAM
The whole model was pinned to reserved RAM: 11 large blocks spread across 2706.91 MB
Hooked to model 'transformer' (WanModel)
Pytorch compilation of 'transformer' is scheduled for these modules : ['blocks.*'].
Pinning data of 'text_encoder' to reserved RAM
The whole model was pinned to reserved RAM: 21 large blocks spread across 6445.29 MB
Hooked to model 'text_encoder' (T5Encoder)
Pinning data of 'vae' to reserved RAM
The whole model was pinned to reserved RAM: 1 large blocks spread across 242.03 MB
Hooked to model 'vae' (WanVAE_)
0%| 0/30 [00:00<?, ?steps/s]
```



```
ty6k-prxm100037:~$ nvidia-smi
```


```
Ubuntu@0360-kcl2-ty6k-prxm100037: ~
Device 0 [NVIDIA RTX 6000 Ada Generation] PCIe GEN 4@16x RX: 1.000 KiB/s TX: 0.000 KiB/s
GPU 2505MHz MEM 9500MHz TEMP 42°C FAN 30% POW 62 / 300 W
GPU[ 0%] MEM[||||] 5.226Gi/47.988Gi]

100 GPU0 %
GPU0 mem%
75
50
25
0


PID USER DEV TYPE GPU GPU MEM CPU HOST MEM Command
1091249 N/A 0 Compute 0% 4796MiB 10% 31% 13046MiB python3 wgp.py --listen --profile 1 --atte
1193 root 0 Graphic 0% 10MiB 0% 0% 12MiB /usr/lib/xorg/Xorg vt1 -displayfd 3 -auth
2022 gdm 0 Graphic 0% 8MiB 0% 1% 48MiB /usr/bin/gnome-shell



F2Setup F6Sort F9Kill F10Quit F12Save Config
```





It generated a short 5 seconds video of the text prompt provided






```

Hooked to model 'transformer' (WanModel)
Pytorch compilation of 'transformer' is scheduled for these modules : ['blocks.*'].
Pinning data of 'text_encoder' to reserved RAM
The whole model was pinned to reserved RAM: 21 large blocks spread across 6445.29 MB
Hooked to model 'text_encoder' (T5Encoder)
Pinning data of 'vae' to reserved RAM
The whole model was pinned to reserved RAM: 1 large blocks spread across 242.03 MB
Hooked to model 'vae' (WanVAE_)
0%|          | 0/30 [00:00<?, ?steps/s]
W1006 23:55:42.734000 35 torch/_dynamo/convert_frame.py:906] [8/256] torch._dynamo hit config.accumulated_
cache size limit (256)
W1006 23:55:42.734000 35 torch/_dynamo/convert_frame.py:906] [8/256] function: 'torch_dynamo_resume_in_
_mm_wrap_torch_nn_modules_linear_Linear_at_6' (<string>:6)
W1006 23:55:42.734000 35 torch/_dynamo/convert_frame.py:906] [8/256] last reason: 8/0: __check_obj_id(
@L['module'].mm.forward.func.__closure__[2].cell_contents, 135977615004160)
W1006 23:55:42.734000 35 torch/_dynamo/convert_frame.py:906] [8/256] To log all recompilation reasons, use
TORCH LOGS="recompiles".
W1006 23:55:42.734000 35 torch/_dynamo/convert_frame.py:906] [8/256] To diagnose recompilation issues, see
https://pytorch.org/docs/main/torch.compiler_troubleshooting.html.
100%|          | 30/30 [03:52<00:00, 7.76s/steps]
New video saved to Path: outputs/2025-10-06-23h57m55s_seed35140621_a baby kangaroo going into the pouch of
mother kangaroo.mp4
/usr/local/lib/python3.10/dist-packages/pygame/pkgdata.py:25: UserWarning: pkg_resources is deprecated as
an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated f
or removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
from pkg_resources import resource_stream, resource_exists
hunyuan_video_720_quantto_int8.safetensor(...): 23%|          | 3.02G/13.2G [00:08<00:20, 488MB/s]

```

Is starts downloading the selected model again



Probably the best text 2 video model available.

Attention mode **sage**, Pytorch compilation **ON**, Data Type **BF16**, Quantization **Scaled Int8**

Choose a Lora Preset or a Settings file in this folder

Apply

Refresh

Save

Delete

Prompts (each new line of prompt will generate a new video, # lines = comments)

a baby kangaroo going into the pouch of mother kangaroo

Category

480p

Resolution Budget (Pixels will be reallocated to preserve Inputs W/H ratio)

832x480 (16:9)

Number of frames (24 frames = 1s), current duration: 3.4s

81

5

337

Number of Inference Steps

30

1

100

☐ Advanced Mode

Set Settings as Default

Export Settings to File

Reset Settings

Load Settings From Video / Image / JSON e Here

Status

Total Generation Time: 8m 27s



Video Info and Late Post Processing & Audio Remuxing

Generate

Queue Management